

Capital One Launchpad Hackathon: Synopsis

1. Team Details

Team Name: Byte Bandits

Team Members:

1. Anshu Sharma
 2. Ambarish Manna
 3. Hardik Kumar
 4. Amitrajeet Konch
-

2. Theme Details

Theme Name: Exploring and Building Agentic AI Solutions for a High-Impact Area of Society: Agriculture – Actionable Hyperlocal Farming Guidance, optimized for real world.

Theme Benefits:

1. Localized Context Aware Assistance: The theme choice enables understanding of local languages, dialects, and cultural contexts, making the system truly accessible and usable by farmers in their native language and communication mode (chat, voice, photo). This hyper localisation is key to the solution.

2. Efficient use of Bandwidth: Reducing data transfer, optimising bandwidth and providing faster response times is crucial in rural areas with intermittent or low-bandwidth internet connectivity. For remote areas, it is paramount that the app work with minimal internet usage.

3. Robust Multi-Modal Input Handling: Ability to process voice, photos, typed queries, and sensor data means sellers can interact in whichever way suits them best. This also allows users to use a different method of communication with the app if one proves ineffective.

4. Integrated Tool Orchestration for Holistic Advice:

The system orchestrates calls to multiple specialized tools like Weather APIs, Market APIs, Soil sensors, and domain-specific knowledge bases (RAG), providing sellers with a 360-degree view of factors impacting their decisions.

5. Seamless Action & Notification Pipeline: Post-hooks and Action Agents automate follow-up steps: sending SMS reminders, scheduling irrigation, placing input orders, or escalating to human agronomists when uncertain.

3. Synopsis

Solution Overview:

Farmora is a smart assistant for farmers that understands natural questions about farming, finds the most relevant data, and gives clear, practical advice all while keeping expensive AI calls to a minimum.

Unlike typical AI chatbots that send your whole question to the cloud immediately, Farmora works offline-first.

It processes, filters, and enriches the question locally, then calls an AI only for the final, human-friendly answer.

The basic crust of Farmora is:

“Data-driven, grounded, relevant and actionable advice to all farm queries”

Farmora takes a user query, uses relevant tools in its toolset to get the most relevant data to answer that query and based on that data, provides answers in the user’s language.

Components & Responsibilities

1. Farmer Client -

Platforms: Android (primary), web, basic feature phone (IVR/SMS).

Inputs: typed question, voice note, photo(s), simple forms (e.g., record harvest).

Local cache for profile, recent advisories, fallback messages.

2. Edge Preprocessor (on device or nearby server) -

Functions: language detection, speech→text, image compression, quick profanity filter.

Tech: small CPU models (fastText / whisper-small / on-device TTS/STT).

Output: normalized JSON (user_id, locale, raw_text, attachments, geo, timestamp).

3. Local Classifier / Router (stateless microservice) -

Multi-label intent detection: weather, market, pest, scheme, irrigation, accounting, order.

Confidence score; if low, route to LLM-based fallback or ask clarifying question.

Tech: DistilBERT / MiniLM / lightweight fine-tuned classifier.

4. Tool Orchestrator -

Calls relevant external/internal tools based on classifier:

Weather service (regional forecasts + alerts).

Market price DB (regional mandi).

Vector DB for RAG (PDFs, scheme docs, agronomy guides).

Pest/disease vision service (local model + RAG).

Soil & IoT sensor API (where available).

Enforces rate limits and caching.

5. Middleware Hooks (pluggable) -

Pre-tool hooks: profanity, PII scrub, quota check.

Tool output validation: schema checks, anomaly detection, confidence gating.

Context compression: token budget enforcement, pruning rules.

Term mapping: local names ↔ canonical names (bidirectional).

Post-LLM hooks: fact-check numbers, safety rules (no illegal pesticides), tone/localization.

6. Context Generator -

Builds minimal, structured prompt: variables + short bullet fact

Keeps fixed system prompt (template) + variable block with compact JSON -> human-readable bulleted summary.

Example structure (max ~400 tokens typical):

SYSTEM: You are a helpful agricultural advisor. Use only supplied context. No hallucinations.

CONTEXT:

- Farmer: Ramesh (ID: 123), Location: Mandya, Karnataka (12.5N,76.9E)

- Crops: Maize (main), Paddy (secondary)

- Stage: Maize - V6

- Weather (next 3 days): Rain 60% day3, Tmax 33C...

- Soil moisture: 28% (sensor id)

- Market: Local mandi price maize 1700₹/q

- RAG: ICAR doc: "blight treatment: apply X" (ref id: vdb:345)

USER QUERY: "Should I irrigate maize tomorrow?"

7. LLM (single reasoning call per query) -

Purpose: synthesize, reason, and produce farmer-facing actions/advice.

System instructions: "Only use supplied context. If missing info, say what you need. Provide 1–3 concrete actions, with rationale and confidence score (0–1)."

Keep the LLM call to one per query where possible.

8. Post Hooks & Action Agents -

Post-checks: numeric sanity, banned recommendations, pesticide safety check.

Action Agents (event-driven microservices): schedule reminders, send SMS/IVR, place orders via dealer APIs.

Escalation: produce "uncertain" flag to forward to human expert.

9. Storage & Indexes -

Farmer profile DB (encrypted): profile, farm geometry, equipment, crop calendar.

Vector DB for RAG: hashed docs, embeddings, metadata.

Time-series DB: market prices, weather history, sensor readings.

Audit logs: all prompts, responses, tool outputs for debugging & compliance.

Technical Stack:

Mobile App: Java, Kotlin, Jetpack Compose, Retrofit/ Ktor, Appwrite Auth SDKs, Gradle

Server Development: Express JS, Flask, Fast API, Selenium

LLM APIs: Groq (lowest latency) – OpenAI OSS, Gemini

Local Models: Flan t5-small, Bert

Database: ChromaDB (RAG), MongoDB, Redis (Caching)

LLM Framework: Langchain + Ollama

APIs: OpenWeather API, Mosdac by ISRO,

Datasets: AgMarkNet(Crop mandi prices), ICAR Reports (Soil health, crop advisories), PM-KISAN (Finance and scheme linking)

Decision Rationale:

1. Offline-First Architecture

- **Constraint:** Limited access to high-speed internet, especially in remote farming regions.
- **Decision:**
 - **Local Processing** (Edge Preprocessor & Local Classifier/Router) ensures that core functionalities like language detection, speech-to-text, and initial query classification are done offline or at a nearby server. This minimizes dependency on cloud-based AI calls.
 - **Local Cache** on the Farmer Client for profile data, recent advisories, and fallback messages helps in offering a seamless experience even during connectivity drops.

2. Efficient AI Model Usage

- **Constraint:** Cloud-based AI calls can be costly and slow, especially when processing large volumes of user queries.
- **Decision:**
 - **Edge Processing** using light models (fastText/Whisper-small, on-device TTS/STT) reduces the need for frequent, expensive cloud-based calls.
 - **LLM Call Optimization:** Only use a cloud-based LLM (e.g., Groq or OpenAI) for the final, reasoned advice generation, reducing unnecessary calls and improving overall system efficiency.

3. Multi-Label Intent Detection & Classification

- **Constraint:** The system needs to accurately classify and route the queries to the right resources while handling ambiguity or low-confidence queries.
- **Decision:**
 - **Local Classifier:** Using **DistilBERT/MiniLM**, lightweight fine-tuned models are employed to classify user queries into multiple categories like weather, market prices,

pests, schemes, and more. This helps in efficiently directing the query to relevant external/internal tools.

- **Fallback Mechanism:** Low-confidence queries are routed to an LLM-based fallback or flagged for clarification, minimizing errors and ensuring high-quality responses.

4. Action Agents & Task Automation

- **Assumption:** Some queries might require follow-up actions (e.g., reminders, placing orders, or scheduling tasks with agronomists).
- **Constraint:** A certain set of actions need to be triggered automatically based on user queries.
- **Decision:**
 - **Event-Driven Action Agents:** If the query requires scheduling a reminder, sending an SMS, or placing an order with a local dealer, **action agents** trigger those actions automatically. This minimizes the farmer's need to manually track and execute tasks.
 - **Escalation Protocol:** If the system is uncertain about the response or needs further expert validation, it flags the query for escalation to a human expert, ensuring high-quality advice.

5. Scalable and Modular Tool Orchestration

- **Assumption:** The system will need to access a variety of external data sources and tools (weather, market prices, pest/disease detection, IoT sensor data).
- **Constraint:** The tools must be orchestrated efficiently to minimize latency, enforce rate limits, and ensure reliability.
- **Decision:**
 - **Modular Tool Orchestrator:** This component centralizes the decision-making for which external tools to call based on the classifier's output. For example, it calls **weather services**, **market price databases**, and **pest vision services** only when necessary, reducing unnecessary calls to external APIs and speeding up the response process.
 - **Caching:** Local caching mechanisms ensure that frequently requested data (e.g., weather forecasts) doesn't require redundant API calls, thereby saving on costs and reducing response time.

Innovation Highlights:

Farmora is an offline-first, resource-conscious AI agent ecosystem tailored for agricultural communities. It combines hyper-local environmental data, domain-specific knowledge, and multi-step reasoning to deliver actionable, context-aware insights while minimizing reliance on costly API calls and internet connectivity. Unlike generic AI assistants, it's purpose-built to integrate live weather, crop health, and local farming practices into decision-making. This makes it both affordable and practical for real-world rural use.

- Offline-first → works even with spotty internet (big plus for rural areas)
- Token-efficient → minimizes LLM calls by doing most processing locally
- Context-rich → blends real-time weather, soil, crop data with AI reasoning
- Domain-specific → built for agriculture, not general chat
- Agentic → it doesn't just answer, it acts, plans, and monitors over time

Feasibility and User-Friendliness:

1. Realism

Feasibility: The tech stack (vector DB, local preprocessing, minimal LLM calls) is realistic with current open-source tools (FAISS, Ollama, Groq API).

Data availability: Weather APIs (Open-Meteo, WeatherAPI) are free for basic use. Agricultural datasets exist (FAO, government open-data portals).

Deployment context: Since rural connectivity is spotty, an offline-first design is not just realistic, but essential.

Hardware constraints: Running light inference locally is possible on low-power devices with quantized models, though heavier tasks will need server fallback.

Verdict: Highly realistic for a prototype, with careful scope control.

2. User-Friendliness

Adoption barrier: Farmers generally won't care about "agents" or "RAG", they'll care about "I ask, I get a useful answer". Simplicity is key.

Interface design: Is multilingual, possibly voice-first for inclusivity. Mobile-first approach. Reduces barrier of entry for farmers.

Trust: Agricultural advice affects livelihoods meaning the answers need clear, simple explanations ("why" this is the recommendation) to build trust. The solution goes through multi step reasoning and multiple checks and balances to reach the answer.

Learning curve: Any farmer that has been using apps like ChatGPT, Gemini etc already will have no problem adopting Farmora as it has similar interface. Even those who don't will not be completely lost due to the intuitive UI.

Verdict: With the right UX decisions (voice + local language + simple flows), adoption could be strong.

3. Operational Efficiency

Token efficiency: Local classification, preprocessing, and context assembly will cut API costs as most processing will be done on-device.

Speed: Vector DB + lightweight models for first-pass analysis = fast responses even on low-end devices.

Scalability: Works in low-resource settings without constant cloud dependence, so scaling to many users is cost-effective. Also switches to the internet whenever applicable

Verdict: Very efficient because it is designed to offload as much as possible to local compute.

4. Long-Term Success Potential

Social impact: Bridges digital divide in agriculture by providing relevant, timely, and localized help.

Economic sustainability: Token cost savings + offline capability makes it affordable to run at scale.

Growth potential: Can expand into other rural services like market prices, government scheme info, supply chain optimization

Retention: If users see improved yields, they'll keep using it; word-of-mouth adoption in rural communities is strong.

Risk: The biggest threat is data trustworthiness. Poor or wrong advice could ruin credibility fast making strong verification loops an important part of the architecture.

4. Methodology/Architecture Diagram

i. App Architecture

https://drive.google.com/file/d/1ai5e60r7x-9l3450G_iEaA4Y5dOJWuLf/view?usp=sharing

ii. Agent Architecture

<https://drive.google.com/file/d/16iEWR21oW24VXqrAya4RnuizEM0YTpcr/view?usp=sharing>
